

Conditioning as Group Action: Lie-Algebra Feature Modulation Composes Where FiLM and Hypernetworks Memorize

Richard Bao
richardbao419@gmail.com

Abstract

Networks are conditioned by scaling features (FiLM) or by generating weights (hypernetworks). Both learn what each condition does. Neither learns what two conditions do together unless training data shows them together. We condition by rotating features, with rotation angles a linear function of the condition: $y = T(c)x + \beta(c)$ with $T(c) = P R(Wc) P^\top$. Rotations add, so $T(c_1+c_2) = T(c_1)T(c_2)$ holds exactly for any weights, and composition becomes a property of the parametrization instead of a behavior learned from data. We test this under a pre-registered protocol: margins, statistical tests, and data splits committed before every run, conditioning cost capped at $1.2\times$ FiLM, and every verdict reported as the gate scored it. Across five transformation suites, Lie conditioning posts the lowest error on never-trained condition combinations: by up to five orders of magnitude when conditions form a group, and by 42.9–53.8% over the strongest hypernetwork baseline on images ($p=9.1 \times 10^{-5}$, $N=10$), with the fewest parameters of any conditioner. An ablation that swaps only the linear map for an MLP loses by up to 43,000 $\times$: the linear map is the mechanism. More expressive conditioners generalize worse; their capacity memorizes training combinations. Three registered gates returned negatives that bound the method: a 1.46 $\times$ fit penalty on rich scenes; a 7 $\times$ failure when the condition specifies content, because rotations move information and cannot add it; and no rollout advantage without a consistency objective, since representation consistency, not operator composition, limits world models. The negatives specify a successor, validated under two further gates: Complex FiLM, an expressive magnitude channel for content and a linear phase channel for composition, beats FiLM by 36% on unseen combinations and matches it on content (0.97 $\times$). One operator covers both roles at FiLM-class cost.

1 Introduction

Conditioning, injecting a class, an action, a timestep, or a text embedding into a network’s computation, recurs in nearly every conditional model in deep learning. Diffusion models modulate every block on the timestep and class via adaLN [2]; visual reasoning, speech, and RL systems modulate features on questions, speakers, and goals [1]; world models condition latent dynamics on actions. Almost all of this conditioning is carried by three mechanisms: concatenation, feature-wise affine modulation (FiLM), or a hypernetwork that emits weights from the condition.

These mechanisms sit at two extremes of an expressivity axis. FiLM’s diagonal transform is cheap and stable but cannot rotate, mix, or exchange information between features. Hypernetworks can express any linear operator but pay quadratic cost and, as we show, systematically *memorize* the training conditions: their error on unseen combinations of familiar conditions is worse than FiLM’s, despite 43 $\times$ the conditioning compute. The middle of this axis, structured operator families between diagonal and dense, is well explored for *weight* fine-tuning (orthogonal and low-rank adapters: LoRA, OFT, BOFT, HRA [3–6]), but largely unexplored in FiLM’s role: *per-sample, input-conditioned modulation of activations*.

We ask a sharp version of the question: *which structure, at FiLM-level cost, buys compositional generalization: correct behavior on combinations of conditions never seen together in training?* Our answer imports a mechanism from positional encoding. RoPE [10] and its generalization GRAPE [11] encode position n as a group action $G(n) = \exp(n \omega L)$: because position enters *linearly in the Lie algebra*, the relative law $G(n+m) = G(n)G(m)$ holds exactly. We transplant this construction from scalar positions in attention to arbitrary condition vectors in activation modulation:

$$y = T(c)x + \beta(c), \quad T(c) = P R(Wc) P^\top, \quad (1)$$

where $R(\theta)$ applies planar 2×2 rotations with angle vector θ , W is a *linear, bias-free* map from the condition into the angle space (the maximal torus of $SO(d)$), and P is a learned orthogonal change of basis. Same-plane rotations commute, so for any W and P ,

$$T(c_1 + c_2) = T(c_1)T(c_2), \quad (2)$$

exactly, at every point in training. Composition of conditions is not learned; it is a theorem about the parametrization. The operator is exactly orthogonal (stable spectra), is the identity at initialization and at $c = 0$, and costs $1.11 \times$ FiLM’s conditioning FLOPs when P is parametrized as a structured (Group-and-Shuffle style [7]) orthogonal; a dense learned basis provably cannot meet this budget (§3.3).

Contributions. (1) We formulate conditioning as a spectrum of per-sample structured operators generalizing FiLM, positioned explicitly against the PEFT literature that owns these operator families for *weight* adaptation (§2). (2) We introduce Lie conditioning (Eq. 1), which makes condition composition exact by construction at FiLM-class cost, including a budget analysis showing dense learned bases exceed any FiLM-comparable FLOP ceiling while structured orthogonal bases fit (§3). (3) We contribute a pre-registered, budget-fair evaluation protocol: margins, statistical tests, and single-read OOD splits fixed before every run; equal-parameter, equal-FLOP comparisons enforced by a shared counter; we release it end-to-end (§4). (4) Across five benchmarks we show: near machine-precision compositional generalization and zero-shot extrapolation to longer compositions when conditions form a group; a reliable $\geq 2 \times$ advantage on real-image transformation with learned latents, surviving categorical (non-group) factors; a mechanism-isolating ablation; interpretable learned rotations that recover ground-truth generators to 2×10^{-4} rad; and a consistent *capacity inversion*: the more expressive the conditioner, the worse its compositional OOD (§5).

2 Related work

FiLM and conditional normalization. FiLM [1] and its conditional normalization variants (conditional BN/LN, adaLN in DiT [2]) apply per-feature affine transforms generated from a condition. All are diagonal: no feature mixing. We keep FiLM’s role: per-sample activation modulation: and enrich only the operator.

Structured operators in PEFT. The operator families between diagonal and dense are well studied as *weight* fine-tuning: low-rank (LoRA [3]), block-diagonal orthogonal (OFT [4]), butterfly orthogonal (BOFT [5]), Householder (HRA [6]), block-orthogonal-with-permutations (Group-and-Shuffle [7]); He et al. [8] give a unified view. These learn a *fixed* transform of weights per task. We claim no novelty for the operator families themselves; our setting differs on every axis that matters: per-sample, input-conditioned, activation-space, and evaluated for compositional generation rather than task adaptation. Hypernetwork-generated adapters (e.g., input-conditioned LoRA variants [9]) generate weights from conditions but inherit the memorization behavior we quantify in §5.

Position encodings as group actions. RoPE [10] rotates query/key pairs by angles linear in position; GRAPE [11] generalizes this to group representations $G(n) = \exp(n\omega L)$ with exact relative laws. The mechanism, argument linear in a Lie algebra, hence exact composition, is theirs; our contribution is transplanting it from scalar position in attention to arbitrary condition vectors in FiLM’s role, learning the basis P , and demonstrating budget-fair compositional gains there.

Compositional generalization in generative models. Montero et al. [12] showed that reconstruction-trained generative models fail to render unseen factor combinations on dSprites-class data, largely independent of disentanglement. Our image experiments adopt that recombination protocol but move the question to the *conditioning mechanism*: with backbone, data, and training identical, which operator family recombines? Symmetry-based representation learning [13, 14] argues latents should carry group actions; we provide a conditioning-side mechanism and controlled evidence for it.

3 Method

3.1 Conditioning as a structured operator

FiLM is the diagonal case $T(c) = \text{diag}(1 + \gamma(c))$ of $y = T(c)x + \beta(c)$. We require of any richer $T(c)$: **(i) identity at init** ($T \equiv I, \beta \equiv 0$), so the conditioned model starts as the unconditioned one; **(ii) bounded spectrum** ($\sigma_{\max}(T(c)) \approx 1$), so modulation cannot blow up activations; **(iii) compositionality**: ideally $T(c_1 + c_2) = T(c_1)T(c_2)$; **(iv) FiLM-class cost**. All four are unit-tested properties of our released implementation.

3.2 Lie conditioning: the condition enters linearly in the Lie algebra

In Eq. 1, $W \in \mathbb{R}^{(d/2) \times k}$ is linear with no bias, so $c = 0$ gives $T = I$, and $W=0$ at initialization gives identity-init. $R(\theta)$ rotates coordinate pairs (x_{2i}, x_{2i+1}) by θ_i : the maximal torus of $\text{SO}(d)$, applied in closed form (no matrix exponential at runtime).

Proposition 1. *For any W, P and any c_1, c_2 : $T(c_1 + c_2) = T(c_1)T(c_2)$, and $T(c)^\top T(c) = I$.*

Proof. $R(\theta_1)R(\theta_2) = R(\theta_1 + \theta_2)$ since coordinate-pair rotations act on disjoint planes; W linear gives $W(c_1 + c_2) = Wc_1 + Wc_2$; conjugation by a fixed P preserves both products. Orthogonality: R and P are orthogonal. $\square$

The point of the construction is what it removes: an MLP head from c to angles (or to any operator parameters) must *learn* composition from data and extrapolate it to unseen combinations. Here composition holds at initialization, during training, and off-distribution. The learned parts, W ’s columns (one generator per condition dimension) and the basis P , only determine *which* planes and *how fast* each condition dimension rotates them.

3.3 The basis must be structured: a budget analysis

$T(c)$ must act in the right basis: a coordinate-aligned torus fails whenever the data’s factors mix coordinates (S2 below shows this collapse). A *dense* learned orthogonal P costs two $d \times d$ matrix–vector products per sample: $4d^2$ FLOPs, i.e. $1.52 \times$ FiLM’s entire conditioning budget at $d=128$, before counting anything else. (An early version of our own accounting under-counted this by $2 \times$; the corrected counter is a regression-tested part of the release, and the erratum is documented in the repository.) We therefore parametrize P as $L=5$ layers of block-diagonal orthogonal matrices (eight 16×16 Cayley blocks per layer, $Q = (I+A)^{-1}(I-A)$ for skew A , exactly orthogonal, identity at $A=0$) with fixed permutations between layers, in the spirit of BOFT and Group-and-Shuffle [5, 7]. Total conditioning cost: $1.11 \times$ FiLM FLOPs, with P undercomplete (4,800 parameters vs. $\dim \text{SO}(128) = 8,128$): it need only discover the planes the conditions act on, not an arbitrary rotation.

4 Pre-registered, budget-fair protocol

Arms. Six conditioners share one condition encoder and differ only in the operator: FiLM; Concat-MLP; conditional LayerNorm; **Hypernetwork** ($T = I + \Delta W(c)$, dense); **Dynamic low-rank** ($T = I + A(c)B(c)^\top$, per-sample LoRA-style); and Lie conditioning. A seventh reported (never gated) arm isolates the mechanism: identical structured P , but an MLP angle head replacing W : *more* compute than ours, differing only in linearity.

Fairness. A single shared counter measures trainable parameters and per-sample forward FLOPs of the whole conditioning path. Gates: ours must hold $\leq 1.05 \times$ the smallest hypernetwork baseline’s parameters and $\leq 1.20 \times$ FiLM’s FLOPs; unstructured baselines may be arbitrarily larger (they are: up to $48 \times$ our parameters). Table 2 reports the budgets. All arms share identical backbones, optimizers, schedules, seeds, and (for image suites) identical evaluation triplets.

Pre-registration. For every suite, the success margin, significance test, seed count, and split protocol were written down before the run: verdict = win only if ours beats the *strongest hypernetwork-family* baseline (not merely FiLM: beating a diagonal operator on rotation-bearing tasks would be nearly tautological) by $\geq 20\%$ relative OOD error, one-sided Mann–Whitney $p \leq 0.01$ and Cliff’s $|\delta| \geq 0.474$ over $N=10$ seeds, at fair budget, with no in-distribution regression. Model selection uses a validation split; each OOD test split is evaluated once per trained arm. Deviations

Table 1: Compositional OOD error (never-trained condition combinations; mean over 10 seeds; lower is better). S1/S2: MSE on $y = M(c)x$. S3–S4: pixel MSE against ground-truth transformed images. Ours is lowest in every column, with complete seed separation vs. the strongest hypernetwork baseline throughout (one-sided Mann–Whitney $p=9.1 \times 10^{-5}$, Cliff’s $\delta=-1.0$). S1–S3b passed the full pre-registered gate; S4’s gate returned a registered negative on a *separate* in-distribution criterion (§5): its OOD cells still favor ours.

Conditioner	S1	S2	S3	S3b	S4	S7
FiLM	1.3×10^{-2}	1.6×10^{-2}	3.0×10^{-3}	5.9×10^{-3}	4.5×10^{-3}	3.1×10^{-3}
Concat-MLP	1.5×10^{-2}	1.1×10^{-2}	2.9×10^{-3}	6.3×10^{-3}	5.4×10^{-3}	–
Cond-LN	2.4×10^{-2}	2.7×10^{-2}	3.2×10^{-3}	6.1×10^{-3}	8.6×10^{-3}	–
Hypernetwork	1.9×10^{-3}	3.4×10^{-3}	3.8×10^{-3}	7.0×10^{-3}	9.3×10^{-3}	3.8×10^{-3}
Dyn. low-rank	5.7×10^{-3}	6.1×10^{-3}	5.1×10^{-3}	8.1×10^{-3}	2.0×10^{-2}	–
MLP-head (abl.)	–	1.1×10^{-3}	2.7×10^{-3}	5.0×10^{-3}	3.3×10^{-3}	–
Complex FiLM (lin.)	–	–	–	–	–	2.0×10^{-3}
Complex FiLM	–	–	–	–	–	2.0×10^{-3}
Lie (ours)	7.4×10^{-4}	3.3×10^{-8}	1.8×10^{-3}	4.0×10^{-3}	1.2×10^{-3}	1.8×10^{-3}

Table 2: Conditioning-path budgets (S3 instantiation). Ours is the cheapest arm.

Conditioner	Params	FLOPs vs. FiLM
FiLM	50,176	1.00×
Concat-MLP	83,072	1.66×
Cond-LN	50,176	1.01×
Hypernetwork	2,147,200	43.17×
Dyn. low-rank	297,856	5.98×
MLP-head (abl.)	52,416	1.27×
Lie (ours)	44,160	1.11×

Table 3: Compositional gap on image suites: OOD/in-dist error ratio (1.0 = no gap). On S3/S3b in-dist error is statistically indistinguishable across arms, isolating composition; on S4 fit quality varies (Concat-MLP and Cond-LN under-fit, deflating their ratios): see §5.

Conditioner	S3	S3b	S4	S7
FiLM	2.08×	2.41×	5.08×	2.07×
Concat-MLP	1.74×	2.45×	1.91×	–
Cond-LN	2.23×	2.56×	1.68×	–
Hypernetwork	2.68×	2.87×	13.11×	2.75×
Dyn. low-rank	3.33×	3.23×	26.10×	–
MLP-head (abl.)	1.89×	2.01×	4.36×	–
Complex FiLM (lin.)	–	–	–	1.42×
Complex FiLM	–	–	–	1.41×
Lie (ours)	1.25×	1.63×	1.13×	1.25×

(two pre-run amendments; one accounting erratum) are disclosed in the repository alongside append-only logs sufficient to recompute every statistic.

Benchmarks. **S1 (synthetic, aligned):** inputs $x \in \mathbb{R}^{128}$; a multi-hot condition over $K=8$ primitive block-plane rotations requests $y = M(c)x$; train on singletons + 8 pairs, test on 10 held-out pairs. **S2 (synthetic, hidden basis):** identical, but $M(c) = BR(c)B^\top$ for a fixed random orthonormal B : the operator is *not* handed the data’s basis and must discover it; we additionally evaluate all 56 never-trained *triple* compositions. **S3 (dSprites):** conditional image transformation. Encode a real image x_1 , apply the conditioning arm for a factor-change Δ over {scale, orientation, position_x, position_y}, decode, and score pixel MSE against the *ground-truth* image at the shifted factors (dSprites [15] is a deterministic factor→image map, so compositional generation is objectively scoreable). Train on single-factor changes + 2 of 6 two-factor types; OOD-test on 2 never-trained two-factor types; report never-trained three-factor types. **S3b (dSprites + categorical):** adds shape change: a factor that is *not* a one-parameter group: encoded as a one-hot difference (which still adds). **S4 (3D Shapes):** the same protocol on RGB, 3D-rendered scenes [16] with six factors (three cyclic hues, scale, orientation, and categorical shape), $\Delta \in \mathbb{R}^9$.




Figure 1: Compositional OOD error by conditioner across suites (log scale; bars: mean $\pm$ sd over 10 seeds; blue = ours, orange = mechanism ablation). Lie conditioning is lowest in every suite: by five orders of magnitude when conditions form a group (S2), and by a 42.9–53.8% margin on real-image suites with learned latents.




Figure 2: **Left (S2):** extrapolation to composition length 3: conditions strictly longer than anything trained. Ours is flat at $\sim 10^{-8}$; every baseline, including the identical-basis MLP-head ablation, degrades. **Right (S1):** learned singleton angles recover the true generators to 2×10^{-4} rad: the operator is directly interpretable.

5 Results

When conditions form a group, composition is solved, not approximated (S1–S2). On the hidden-basis suite, ours reaches 3.3×10^{-8} OOD MSE, five orders of magnitude below the strongest baseline (Table 1), and, on the 56 never-trained *triples*, 5.3×10^{-8} : error does not grow with composition length (Fig. 2, left). The composition identity (Eq. 2) holds on trained models to 6.7×10^{-6} (Frobenius). A coordinate-aligned ablation ($P=I$) collapses to FiLM’s error: basis discovery, not hand-alignment, carries the result: and the learned angles match the ground-truth generators to 2×10^{-4} rad (Fig. 2, right): one interpretable rotation plane per primitive.

The advantage survives learned latents and real images (S3–S4). On dSprites transformation, all arms fit in-distribution identically ($\approx 1.4 \times 10^{-3}$ pixel MSE), the comparison isolates composition by construction, yet ours is the only arm whose OOD/in-dist ratio stays near 1 ($1.25\times$ vs. $2.1\text{--}3.3\times$ for all baselines; Table 3), beating the strongest baseline by 53.8% at $39\times$ fewer conditioning FLOPs. Never-trained three-factor changes degrade ours most gracefully ($1.48\times$ its own in-dist).

Categorical factors attenuate but do not break the advantage (S3b). Adding shape change: a factor with no one-parameter group structure, encoded as a one-hot difference: attenuates the margin exactly as the theory predicts (exact composition no longer covers every factor), yet the pre-registered gate still passes decisively: 42.9% over the strongest baseline ($p=9.1 \times 10^{-5}$, complete seed separation), with a compositional gap of $1.63\times$ vs. $2.4\text{--}3.2\times$ for every baseline (Table 3). The linear map simply assigns each categorical transition its own rotation, sacrificing exactness on the

categorical dimension while preserving it on the geometric ones. The capacity inversion persists here too (hypernetwork 2.87×, dynamic low-rank 3.23×, both worse than FiLM’s 2.41×).

S4 (3D Shapes): the gate fires a registered negative: and quantifies the trade-off. On RGB, 3D-rendered scenes with six factors, ours again has the best compositional metrics by a wide margin: OOD error 87.5% below the strongest baseline (complete seed separation), OOD/in-dist gap 1.13× vs. 13.1× (hypernetwork) and 26.1× (dynamic low-rank), and the flattest degradation on never-trained three-factor changes. Nevertheless, the pre-registered verdict for this suite is *negative*: our in-distribution error (1.03×10^{-3}) is 1.46× the best memorizer’s (7.1×10^{-4}), violating the registered no-regression criterion ($\leq 1.10\times$) that exists precisely to prevent buying OOD robustness by under-fitting. We report the verdict as registered. The reading we defend: on richer scenes the exact-orthogonality constraint has a measurable in-distribution price, and whether that price is worth an 8× OOD improvement is application-dependent: a trade-off our protocol surfaces with numbers where a less constrained evaluation would have simply declared victory. (Two baselines’ small ratios in Table 3 are artifacts of poor fits inflating the denominator: Concat-MLP and Cond-LN fit 4–7× worse in-distribution than the memorizers.)¹

S5 (content role): rotation cannot inject content. A width-128 diffusion transformer trained factor-conditionally on dSprites (full conditioning determines the target image uniquely, so generation on held-out combinations is scored by pixel MSE against ground truth) yields the second registered negative: the rotation class path underfits at 0.296 vs. adaLN’s 0.042, and the MLP-head ablation fails identically (0.308). The failure therefore sits in the channel, not the linearity: a rotation preserves information, so it transforms representations well and injects content poorly. Every suite the method wins conditions on *changes*; this suite conditions on *content*. The distinction gives practitioners a usage rule and gives the framework its next design point (a magnitude channel; §7).

S6 (rollout role): consistency, not composition, limits world models. An action-conditioned world model applies the conditioning arm recurrently as the latent transition; horizons of 10 and 20 actions (training capped at 3) probe error compounding. Every arm compounds identically (all reach ≈ 0.043 at length 20; our margin over the strongest baseline is 7%, far below the registered 30%), and the mildly contractive dynamic low-rank arm degrades least (0.015 at length 20). The mechanism: each step introduces a small mismatch between $T(a)z(s)$ and $z(s+1)$; a contractive transition reabsorbs that mismatch toward the decodable manifold at every step, while an isometry preserves it indefinitely. Exact operator composition cannot repair an inconsistent representation. This falsified our own pre-registered prediction, and we report it as registered. A follow-up suite adds the missing piece: training with a latent-consistency loss ($\lambda \|T(a)z_s - z_{s+1}\|^2$, applied to every arm equally) inverts the picture. The Lie arm’s rollout error goes flat, 0.0039 at horizon 20 (growth 2.0×) against the hypernetwork’s 0.0384 (growth 19.8×) and FiLM’s 0.0173: a 90% margin. That suite’s registered gate still returns KILL, because its second criterion, a single-step pairs margin, fails: consistency training equalizes single steps across arms. We report both facts. The consistency loss is standard practice in model-based RL [19], and rotation transitions with latent-prediction losses appear in the homomorphism autoencoder [18]; what is new here is the controlled decomposition at matched budget: the loss flattens rollouts for isometric operators specifically, both rotation arms achieve it, and the linear map’s contribution washes out at this site. A further registered suite rules out the cheap alternative: fixed uniform contraction at rates from 0.003 to 0.1 leaves horizon-20 error unchanged, so the correction must be adaptive rather than a constant.

S7: the negatives specify the successor, and it passes both gates. Complex FiLM multiplies each feature pair by a complex number: magnitude from FiLM’s expressive MLP head (the content channel), phase linear in the condition (the composition channel). Two pre-registered gates test the two roles. In the S3 transformation harness it beats FiLM by 36.1% on unseen change types at fit parity and 0.84× FiLM’s conditioning cost. In the S5 diffusion harness it generates held-out combinations at 0.97× FiLM’s error, restoring the content ability the pure rotation lacked by 7×. A linear magnitude channel fails (error 0.19 vs. film’s 0.04): content needs the expressive head, exactly as the boundary map predicted. FiLM is the real part of the working operator; the whole complex number does both jobs.

S8: guidance as an operator power. CFG strengthens conditioning by linear extrapolation in noise space, and its distortion at high scales is established [24]. Because dSprites conditioning is deterministic, our harness isolates that

¹All suite cells are generated from the same pre-registered pipeline; per-suite verdicts and append-only logs are in the repository.

distortion objectively against an exact alternative: with Complex FiLM, feeding αc powers the phase operator exactly, with no unconditional pass and no distillation [25]. Under three pre-registered gates the comparison confirms: parity at native strength ($0.96\times$ FiLM+CFG), error growth of $23.7\times$ from strength 1 to 8 against CFG’s $50.7\times$ ($p=1.6\times 10^{-4}$), and $2\text{--}5\times$ lower error at every strength above 1. Neither method improves over strength 1 on this instrument; the target is deterministic, so guidance can only distort, as the suite’s registered caveat stated in advance.

The mechanism is the linearity, and capacity actively hurts. The ablation shares our orthogonal basis and exceeds our compute; it differs *only* in replacing the linear angle map with an MLP. It is $43,000\times$ worse on synthetic triples and $1.55\times$ worse on dSprites OOD. Meanwhile the two hypernetwork-family arms, the only arms that can represent *any* conditioning operator, have the *worst* compositional gaps on images ($2.7\times$ and $3.3\times$) despite $6\text{--}43\times$ the FLOPs: they spend capacity memorizing training condition-types. Compositional conditioning is an inductive-bias problem; adding expressivity moves it in the wrong direction.

6 Practical guidance

For practitioners, the eight suites reduce to three plain rules. **(1)** If your conditions are changes that combine (edit deltas, attribute offsets, pose or camera adjustments, applied once), use the Lie operator: same cost class as FiLM, fewer parameters, identical in-distribution quality, and far better behavior on combinations your training data never contained; FiLM’s failure on such combinations is silent, so the structural guarantee has operational value. **(2)** If your conditions specify what to generate (class labels, text, identity), keep FiLM/adaLN: our own registered negative shows the rotation channel is $7\times$ worse there, because rotations move information and cannot add it. **(3)** In recurrent rollouts, pair an isometric operator with a latent-consistency loss; without the loss every conditioner compounded error at the same rate, and with it the rotation operator’s rollouts stay flat (90% below the strongest baseline at horizon 20). All evidence is at controlled 64×64 scale; the adaLN-site variant is cheaper than FiLM’s marginal class path, so trying the swap is inexpensive where rule (1) applies.

7 Limitations and scope

Learned latents do not support exact group action: the real-image advantage is a reliable $\approx 2\times$, not the synthetic suites’ $10^5\times$, and Eq. 2 holding in the operator does not force the encoder–decoder pair to respect it. S4 shows the constraint can carry a real in-distribution cost ($1.46\times$ the best memorizer) on richer scenes, our own gate scores that suite as a negative, so exact orthogonality should be understood as a robustness–fit trade-off, not a free lunch. S5 and S6 add two harder boundaries: the operator conditions on changes, not content, and its guarantees end where representation consistency begins. All three point at the same successor design: the abelian algebra torus $\times$ diagonal, i.e. complex-valued FiLM with magnitude as the content channel and phase as the composition channel, paired with a consistency objective for recurrent use. Our images are controlled 64×64 benchmarks with known factors; natural-image, text-conditioned, and diffusion-scale evaluations (the natural deployment is the adaLN site in DiT blocks) are future work, as is conditioning at multiple depths. Conditions whose semantics are far from any group action, such as open-ended text, unordered categories at large cardinality: may need hybrid operators (orthogonal $\times$ low-rank), which our framework accommodates but this paper does not evaluate at scale.

8 Reproducibility statement

Every experiment was pre-registered (margins, tests, splits fixed before running; specifications in the repository), runs from a single command per suite, and writes append-only logs from which every number in this paper: all tables and figures are generated directly from those logs, can be recomputed. Property tests pin the operator invariants (identity-init, orthogonality, bounded spectrum, exact composition) and the fairness counters. We additionally disclose two pre-run amendments (training loss on sparse binary images; step count) and one corrected FLOP accounting erratum, with the corrected counter regression-tested.

References

- [1] E. Perez, F. Strub, H. de Vries, V. Dumoulin, A. Courville. FiLM: Visual reasoning with a general conditioning layer. *AAAI*, 2018.
- [2] W. Peebles, S. Xie. Scalable diffusion models with transformers. *ICCV*, 2023.
- [3] E. Hu, Y. Shen, P. Wallis, Z. Allen-Zhu, Y. Li, S. Wang, L. Wang, W. Chen. LoRA: Low-rank adaptation of large language models. *ICLR*, 2022.
- [4] Z. Qiu, W. Liu, H. Feng, Y. Xue, Y. Feng, Z. Liu, D. Zhang, A. Weller, B. Schölkopf. Controlling text-to-image diffusion by orthogonal finetuning. *NeurIPS*, 2023.
- [5] W. Liu, Z. Qiu, Y. Feng, Y. Xiu, Y. Xue, L. Yu, H. Feng, Z. Liu, J. Heo, S. Peng, et al. Parameter-efficient orthogonal finetuning via butterfly factorization. *ICLR*, 2024.
- [6] S. Yuan, H. Liu, H. Xu. Bridging the gap between low-rank and orthogonal adaptation via Householder reflection adaptation. *NeurIPS*, 2024.
- [7] M. Gorbunov, N. Yudin, V. Soboleva, A. Alanov, A. Naumov, M. Rakhuba. Group and shuffle: Efficient structured orthogonal parametrization. *NeurIPS*, 2024.
- [8] J. He, C. Zhou, X. Ma, T. Berg-Kirkpatrick, G. Neubig. Towards a unified view of parameter-efficient transfer learning. *ICLR*, 2022.
- [9] D. Ha, A. Dai, Q. Le. Hypernetworks. *ICLR*, 2017.
- [10] J. Su, Y. Lu, S. Pan, A. Murtadha, B. Wen, Y. Liu. RoFormer: Enhanced transformer with rotary position embedding. *arXiv:2104.09864*, 2021.
- [11] Anonymous. Group representational position encoding. *arXiv:2512.07805*; *ICLR*, 2026.
- [12] M. Montero, C. Ludwig, R. Costa, G. Malhotra, J. Bowers. The role of disentanglement in generalisation. *ICLR*, 2021.
- [13] I. Higgins, D. Amos, D. Pfau, S. Racaniere, L. Matthey, D. Rezende, A. Lerchner. Towards a definition of disentangled representations. *arXiv:1812.02230*, 2018.
- [14] T. Cohen, M. Welling. Group equivariant convolutional networks. *ICML*, 2016.
- [15] L. Matthey, I. Higgins, D. Hassabis, A. Lerchner. dSprites: Disentanglement testing sprites dataset. <https://github.com/deepmind/dsprites-dataset>, 2017.
- [16] C. Burgess, H. Kim. 3D Shapes dataset. <https://github.com/deepmind/3d-shapes>, 2018.
- [17] D. Worrall, S. Garbin, D. Turmukhambetov, G. Brostow. Interpretable transformations with encoder-decoder networks. *ICCV*, 2017.
- [18] H. Keurti, A. Neitz, S. Weichwald, B. Schölkopf, et al. Homomorphism autoencoder: Learning group structured representations from observed transitions. *ICML*, 2023.
- [19] N. Hansen, X. Wang, H. Su. Temporal difference learning for model predictive control. *ICML*, 2022.
- [20] E. Chan, M. Monteiro, P. Kellnhofer, J. Wu, G. Wetzstein. pi-GAN: Periodic implicit generative adversarial networks. *CVPR*, 2021.
- [21] Z. Sun, Z. Deng, J. Nie, J. Tang. RotatE: Knowledge graph embedding by relational rotation in complex space. *ICLR*, 2019.
- [22] X. Zhu, C. Xu, D. Tao. Commutative Lie group VAE for disentanglement learning. *ICML*, 2021.

- [23] N. Liu, S. Li, Y. Du, A. Torralba, J. Tenenbaum. Compositional visual generation with composable diffusion models. *ECCV*, 2022.
- [24] A. Bradley, P. Nakkiran. Classifier-free guidance is a predictor-corrector. *arXiv:2408.09000*, 2024.
- [25] M. Fu et al. TeEFusion: Blending text embeddings to distill classifier-free guidance. *ICCV*, 2025.
- [26] T. Nagarajan, K. Grauman. Attributes as operators. *ECCV*, 2018.
- [27] M. Georgopoulos, G. Chrysos, M. Pantic, Y. Panagakis. Multilinear latent conditioning for generating unseen attribute combinations. *ICML*, 2020.