

# CGA: Conditioning as Group Action

Richard Bao

richardbao419@gmail.com

## Abstract

Networks are conditioned by scaling features (FiLM) or by generating weights (hypernetworks). Both learn what each condition does. Neither learns what two conditions do together unless training data shows them together. We condition by rotating features, with rotation angles a linear function of the condition. Rotations add, so combined conditions compose automatically, and the network handles combinations it never saw. **CGA** (Conditioning as Group Action) formalizes this: the condition enters the operator linearly in an abelian Lie algebra,  $T(c) = P \exp(A(Wc))P^\top$ , so  $T(c_1+c_2) = T(c_1)T(c_2)$  holds for any weights. A diagonal algebra yields a compositional FiLM variant; the rotation algebra with scalar position is RoPE. Under a pre-registered protocol (margins, tests, and splits committed before every run; conditioning cost within  $1.2\times$  FiLM), CGA posts the lowest error on never-trained condition combinations across five transformation suites, with complete seed separation ( $p=9.1\times 10^{-5}$ ,  $N=10$ ) and zero-shot extrapolation to combinations longer than any trained. Swapping its linear map for an MLP raises error up to  $4.3\times 10^4$ ; the most expressive baselines generalize worst. Three registered gates returned negatives that bound the method: a fit penalty on rich scenes, a content-conditioning failure (rotations move information and cannot add it), and rollouts limited by representation consistency. The negatives specify the successor. **Complex FiLM**, magnitude for content and phase for composition, beats FiLM by 36% on unseen combinations and matches it on content ( $0.97\times$ ) under two further pre-registered gates: one operator, both roles, FiLM cost.

## 1 The framework

Let  $x \in \mathbb{R}^d$  be features and  $c \in \mathbb{R}^k$  a condition. FiLM [1] computes  $y = \gamma(c) \odot x + \beta(c)$ . AdaLN in DiT blocks [2] is its normalized form. Hypernetworks emit full weights  $W(c)$  [3]. All of these are instances of

$$y = T(c)x + \beta(c), \quad (1)$$

and they differ only in where the per-sample operator  $T(c)$  comes from.

**Definition 1** (Compositional conditioner).  $T : \mathbb{R}^k \rightarrow \text{GL}(d)$  is compositional if it is a homomorphism from  $(\mathbb{R}^k, +)$ :  $T(c_1+c_2) = T(c_1)T(c_2)$  and  $T(0) = I$ .

**Proposition 1** (Exact composition by construction). Let  $\mathfrak{a}$  be an abelian subalgebra of  $\mathfrak{gl}(d)$ ,  $A : \mathbb{R}^m \rightarrow \mathfrak{a}$  linear,  $W \in \mathbb{R}^{m \times k}$ , and  $P$  orthogonal. Then  $T(c) = P \exp(A(Wc))P^\top$  is a compositional conditioner for every  $W, P$ . If  $\mathfrak{a}$  is skew-symmetric,  $T(c)$  is orthogonal.

*Proof.*  $\exp(X)\exp(Y) = \exp(X+Y)$  for commuting  $X, Y$ ;  $A \circ W$  is linear; conjugation preserves products;  $\exp(0) = I$ .  $\square$

The point is simple. Composition is usually a behavior the network must learn from data. Here it is a property of the parametrization. It holds at initialization, during training, and far outside the training distribution. Training only decides *which* directions the condition rotates and *how fast*.

**Special cases.** Choices of  $(\mathfrak{a}, c, P)$  recover deployed mechanisms:

| $\mathfrak{a}$                              | condition                    | $P$        | recovered mechanism                                            |
|---------------------------------------------|------------------------------|------------|----------------------------------------------------------------|
| $\text{diag}(\mathbb{R}^d)$                 | $c$                          | $I$        | <i>exponential FiLM</i> : $y = e^{Wc} \odot x$ (compositional) |
| $\mathfrak{so}(2)^{d/2}$                    | $c = n$ (scalar position)    | $I$        | RoPE [4]                                                       |
| $\mathfrak{so}(2)^{d/2}$                    | $c = n$ , learned planes     | learned    | multiplicative GRAPE [5]                                       |
| $\mathfrak{so}(2)^{d/2}$                    | general $c \in \mathbb{R}^k$ | structured | <b>this work</b> (FiLM’s role)                                 |
| $\text{diag} \times \mathfrak{so}(2)^{d/2}$ | general $c$                  | canonical  | <b>Complex FiLM</b> (the successor; §2)                        |




Figure 1: Error on unseen condition combinations, by conditioner (log scale; blue = CGA, orange = MLP-head ablation).

Standard FiLM reaches the same diagonal operators through an MLP head  $\gamma(c)$ . The operators match; the composition law is missing. Section 2 shows the missing law is exactly what fails on unseen combinations: an ablation with the same operators, the same basis, and more compute loses by up to  $4.3 \times 10^4$  once its map from condition to operator is an MLP. The PEFT literature (LoRA, OFT, BOFT, HRA, Group-and-Shuffle [6–10]) uses these operator families to fine-tune *weights*, one fixed transform per task. CGA uses them in FiLM’s role: a new operator per sample, applied to activations, where the composition question lives.

**Cost.** With  $\alpha = \pi(2)^{d/2}$ ,  $\exp$  is a closed-form planar rotation:  $3d$  FLOPs, no matrix exponential. The basis  $P$  is the whole budget question. A dense learned  $P$  costs  $4d^2$  FLOPs per sample, which is  $1.52\times$  FiLM’s entire conditioning path at  $d=128$ . Two fixes work. First, a structured orthogonal  $P$  (five block-diagonal Cayley layers with fixed shuffles, following Group-and-Shuffle) brings the total to  $1.11\times$  FiLM. This  $P$  has 4,800 parameters against the 8,128 dimensions of  $\text{SO}(128)$ , and that is enough: it only needs to find the planes the condition acts on. Second, inside a transformer no explicit  $P$  is needed at all. The dense projections next to the rotation learn the basis, which is how RoPE works. The operator equals the identity at initialization and at  $c=0$ , and every singular value is 1. Unit tests pin all three properties.

## 2 Experiments

**Protocol.** Every suite is pre-registered. We fixed the success margins, the tests ( $N=10$  seeds, one-sided Mann–Whitney  $p \leq 0.01$ , Cliff’s  $|\delta| \geq 0.474$ ), and the data splits before each run, and committed them to version control. Each test split is read once. All arms share one backbone and one training recipe. A shared counter enforces the budget: our conditioning path stays within  $1.20\times$  FiLM’s FLOPs and  $1.05\times$  the smallest hypernetwork baseline’s parameters. Baselines may be bigger, and reach  $48\times$  our parameter count. We report every verdict, including the negatives. Baselines: FiLM, conditional LayerNorm, Concat-MLP, a dense hypernetwork  $I+\Delta W(c)$ , a dynamic low-rank map  $I+A(c)B(c)^\top$ , and an MLP-head ablation of our own method.

**Suites.** **S1/S2 (synthetic):** learn  $y = M(c)x$ , where  $c$  selects a combination of 8 rotation primitives. S2 hides the primitives behind a random basis and also tests all 56 never-trained *triples*. **S3/S3b (dSprites) and S4 (3D Shapes):** image transformation. Encode a real image, apply  $T(\Delta)$  to the latent for a factor change  $\Delta$ , decode, and score pixel error against the true transformed image. OOD means never-trained *types* of two-factor changes. S3b and S4 add a categorical shape factor. **S5 (content):** a small diffusion transformer generates dSprites images from full factor specifications; held-out combinations have unique ground-truth images, so generation error is objective. **S6 (rollouts):** a world model applies the conditioning arm recurrently, one action per step; we test unseen action-type pairs and rollouts of length 10 and 20 after training on length 3. **S7 (the successor, two gates):** Complex FiLM, magnitude from FiLM’s MLP head and phase linear in the condition, rerun through the S3 transformation suite and the S5 content suite. **S8 (guidance):** both arms train with condition dropout; we sample FiLM with classifier-free guidance at strength  $w$  and Complex FiLM by feeding  $\alpha c$  (an exact operator power on the phase channel), then score generation error against the deterministic ground truth at each strength.

**Findings.** **(1) When conditions form a group, composition is solved.** On S2, CGA reaches  $3.3\times 10^{-8}$  error on unseen pairs and  $5.3\times 10^{-8}$  on never-trained triples. The best baseline sits five orders of magnitude higher, and every baseline gets worse as combinations get longer. CGA does not. Its learned angles match the true generators to  $2\times 10^{-4}$

Table 1: Error on never-trained condition combinations (mean over 10 seeds). CGA is lowest in every column, with complete seed separation against the strongest hypernetwork baseline ( $p=9.1\times 10^{-5}$ ,  $\delta=-1.0$ ). S1–S3b pass the full pre-registered gate; S4 fails a separate fit criterion (text).

| Conditioner         | S1                  | S2                  | S3                  | S3b                 | S4                  | S7                  |
|---------------------|---------------------|---------------------|---------------------|---------------------|---------------------|---------------------|
| FiLM                | $1.3\times 10^{-2}$ | $1.6\times 10^{-2}$ | $3.0\times 10^{-3}$ | $5.9\times 10^{-3}$ | $4.5\times 10^{-3}$ | $3.1\times 10^{-3}$ |
| Concat-MLP          | $1.5\times 10^{-2}$ | $1.1\times 10^{-2}$ | $2.9\times 10^{-3}$ | $6.3\times 10^{-3}$ | $5.4\times 10^{-3}$ | –                   |
| Cond-LN             | $2.4\times 10^{-2}$ | $2.7\times 10^{-2}$ | $3.2\times 10^{-3}$ | $6.1\times 10^{-3}$ | $8.6\times 10^{-3}$ | –                   |
| Hypernetwork        | $1.9\times 10^{-3}$ | $3.4\times 10^{-3}$ | $3.8\times 10^{-3}$ | $7.0\times 10^{-3}$ | $9.3\times 10^{-3}$ | $3.8\times 10^{-3}$ |
| Dyn. low-rank       | $5.7\times 10^{-3}$ | $6.1\times 10^{-3}$ | $5.1\times 10^{-3}$ | $8.1\times 10^{-3}$ | $2.0\times 10^{-2}$ | –                   |
| MLP-head (abl.)     | –                   | $1.1\times 10^{-3}$ | $2.7\times 10^{-3}$ | $5.0\times 10^{-3}$ | $3.3\times 10^{-3}$ | –                   |
| Complex FiLM (lin.) | –                   | –                   | –                   | –                   | –                   | $2.0\times 10^{-3}$ |
| <b>Complex FiLM</b> | –                   | –                   | –                   | –                   | –                   | $2.0\times 10^{-3}$ |
| <b>Lie (ours)</b>   | $7.4\times 10^{-4}$ | $3.3\times 10^{-8}$ | $1.8\times 10^{-3}$ | $4.0\times 10^{-3}$ | $1.2\times 10^{-3}$ | $1.8\times 10^{-3}$ |

rad, so the operator is directly readable. **(2) The advantage survives learned representations.** On dSprites, every arm fits the training distribution equally well, so the comparison isolates composition. CGA’s error grows  $1.25\times$  from training to unseen combinations ( $1.63\times$  with the categorical factor). Every baseline grows  $2.1\text{--}3.3\times$ . CGA beats the strongest baseline by 53.8% (42.9%) at  $39\times$  fewer conditioning FLOPs. **(3) The linear map is the mechanism.** The ablation keeps the basis and the operator family, spends more compute, and swaps only the linear map for an MLP. It loses by  $4.3\times 10^4\times$  on synthetic triples and  $1.55\times$  on dSprites. **(4) More capacity makes it worse.** The two hypernetwork arms can express any operator. They post the worst generalization gaps on images ( $2.7\text{--}26\times$ ). Their capacity goes into memorizing the training combinations. **(5) Three registered negatives.** On 3D Shapes, CGA has the best OOD numbers (87.5% margin) but fits the training distribution  $1.46\times$  worse than the best memorizer, which trips our no-regression gate. In the content suite, rotation conditioning fails (0.31 vs. film’s 0.04), and the ablation fails identically: a rotation moves information around and cannot add any, so it cannot carry “generate this image.” In the rollout suite, every arm compounds error at the same rate, and the mildly contractive low-rank arm degrades least. Each step of a rollout adds a small mismatch between the predicted and true latent. A contraction shrinks that mismatch at every step; an isometry carries it forever. Exact operator composition cannot repair an inconsistent representation. **(6) A consistency loss completes the rollout recipe.** Training with  $\lambda\|T(a)z_s - z_{s+1}\|^2$ , applied to every arm equally, inverts the rollout picture: the rotation operator’s error at horizon 20 goes flat at 0.0039 (growth  $2.0\times$  from its trained range) against the hypernetwork’s 0.038 (growth  $19.8\times$ ) and FiLM’s 0.017. The loss is standard in model-based RL [15], and rotation transitions with latent-prediction objectives appear in the homomorphism autoencoder [14]; our contribution is the controlled decomposition at matched budget, plus a null that sharpens it: fixed uniform contraction at any rate from 0.003 to 0.1 leaves horizon-20 error unchanged, so the correction must be adaptive. That suite’s registered gate still scores KILL on a separate single-step criterion, and we report it so. **(7) The negatives specify the fix, and it works.** Complex FiLM multiplies feature pairs by  $m e^{i\theta}$ : magnitude from FiLM’s expressive head (content), phase linear in the condition (composition). Under two pre-registered gates it beats FiLM by 36% on unseen change types at fit parity and  $0.84\times$  FiLM’s cost, and matches FiLM on content generation at  $0.97\times$  its error, where the pure rotation failed by  $7\times$ . A linear magnitude channel fails (0.19): content needs the expressive head. FiLM is the real part of the working operator; the whole complex number does both jobs. **(8) Guidance without the second pass.** CFG’s distortion at high scales is established [20]; our deterministic instrument isolates it against an exact alternative. CFG’s error grows  $50.7\times$  from strength 1 to 8; powering the condition through Complex FiLM’s phase channel grows  $23.7\times$  ( $p=1.6\times 10^{-4}$ ), stays  $2\text{--}5\times$  lower at every strength above 1, matches at native strength ( $0.96\times$ ), and needs no unconditional pass and no distillation [21]. Neither method improves over strength 1 here: the target is deterministic, so there is no distribution to sharpen, exactly as the suite’s registered caveat stated.

### 3 Practical guidance: when to use this

Three rules for practitioners, from the eight suites.

**Use phase conditioning when your conditions are changes that combine.** Edit deltas, attribute offsets, pose or camera adjustments, applied once. The cost matches FiLM ( $1.11\times$  FLOPs, fewer parameters), training-distribution

quality matches, and behavior on never-trained combinations is 43% to  $10^5\times$  better. FiLM’s failure on new combinations is silent: nothing in training warns you. The guarantee removes that failure class.

**For conditions that say what to generate, use Complex FiLM or keep FiLM.** Pure rotation conditioning fails there ( $7\times$  worse; our registered negative). Complex FiLM matches FiLM on content ( $0.97\times$ ) while keeping the composition wins, so it is the candidate drop-in at our evidence scale; plain FiLM remains the conservative choice.

**In rollouts, pair an isometric operator with a consistency loss.** Without that loss, all conditioners compounded error at the same rate. With it, the rotation operator’s rollout error stays flat where FiLM’s grows  $4.5\times$  and a hypernetwork’s  $10\times$ : train the transition to satisfy  $T(a)z_s \approx z_{s+1}$  and the isometry stops carrying old errors forward because there are none to carry.

All evidence comes from controlled  $64\times 64$  benchmarks with known factors. Nothing here is validated at production scale. The transformer variant costs less than FiLM’s marginal class path, so where the first rule applies, trying the swap is cheap.

## 4 Discussion

CGA does for conditioning what RoPE did for position: it moves composition from a behavior the network must learn into a law the parametrization guarantees. Four boundaries are measured. Categorical factors shrink the advantage. Rich scenes make orthogonality pay a fit penalty. Content specification defeats a transformation channel. Rollouts are limited by representation consistency, which no operator guarantee repairs. The first three pointed at one successor design, the algebra torus  $\times$  diagonal, and S7 validated it: Complex FiLM carries content with its magnitude and composition with its phase. The fourth boundary falls to a consistency objective on the representation: with it, the isometric operator gives flat rollouts where every baseline drifts. Related work has learned group actions on latents for world models [11, 12]; CGA differs by living in FiLM’s role under a matched budget, with composition holding by construction rather than by training objective.

**Reproducibility.** Every suite is pre-registered, and every number in this paper regenerates from committed logs. Unit tests pin the operator invariants and the budget counters. The repository discloses all amendments and one accounting erratum.

## References

- [1] E. Perez et al. FiLM: Visual reasoning with a general conditioning layer. *AAAI*, 2018.
- [2] W. Peebles, S. Xie. Scalable diffusion models with transformers. *ICCV*, 2023.
- [3] D. Ha, A. Dai, Q. Le. Hypernetworks. *ICLR*, 2017.
- [4] J. Su et al. RoFormer: Enhanced transformer with rotary position embedding. *arXiv:2104.09864*, 2021.
- [5] L. Zhang, K. Chen, X. Liu, et al. Group representational position encoding. *ICLR*, 2026. *arXiv:2512.07805*.
- [6] E. Hu et al. LoRA. *ICLR*, 2022.
- [7] Z. Qiu et al. Orthogonal finetuning. *NeurIPS*, 2023.
- [8] W. Liu et al. BOFT. *ICLR*, 2024.
- [9] S. Yuan et al. Householder reflection adaptation. *NeurIPS*, 2024.
- [10] M. Gorbunov et al. Group and shuffle: Efficient structured orthogonal parametrization. *NeurIPS*, 2024.
- [11] R. Quessard, T. Barrett, W. Clements. Learning disentangled representations and group structure of dynamical environments. *NeurIPS*, 2020.
- [12] H. Caselles-Dupré, M. Garcia-Ortiz, D. Filliat. Symmetry-based disentangled representation learning requires interaction with environments. *NeurIPS*, 2019.

- [13] D. Worrall, S. Garbin, D. Turmukhambetov, G. Brostow. Interpretable transformations with encoder-decoder networks. *ICCV*, 2017.
- [14] H. Keurti, A. Neitz, S. Weichwald, B. Schölkopf, et al. Homomorphism autoencoder: Learning group structured representations from observed transitions. *ICML*, 2023.
- [15] N. Hansen, X. Wang, H. Su. Temporal difference learning for model predictive control. *ICML*, 2022.
- [16] E. Chan, M. Monteiro, P. Kellnhofer, J. Wu, G. Wetzstein. pi-GAN: Periodic implicit generative adversarial networks. *CVPR*, 2021.
- [17] Z. Sun, Z. Deng, J. Nie, J. Tang. RotatE: Knowledge graph embedding by relational rotation in complex space. *ICLR*, 2019.
- [18] X. Zhu, C. Xu, D. Tao. Commutative Lie group VAE for disentanglement learning. *ICML*, 2021.
- [19] N. Liu, S. Li, Y. Du, A. Torralba, J. Tenenbaum. Compositional visual generation with composable diffusion models. *ECCV*, 2022.
- [20] A. Bradley, P. Nakkiran. Classifier-free guidance is a predictor-corrector. *arXiv:2408.09000*, 2024.
- [21] M. Fu et al. TeEFusion: Blending text embeddings to distill classifier-free guidance. *ICCV*, 2025.
- [22] T. Nagarajan, K. Grauman. Attributes as operators. *ECCV*, 2018.
- [23] M. Georgopoulos, G. Chrysos, M. Pantic, Y. Panagakis. Multilinear latent conditioning for generating unseen attribute combinations. *ICML*, 2020.